

See in the dark

Deep Learning Homework Paper

Takáts Bálint
Mészáros Gergő
Oszvald Levente

Abstract

Taking pictures in a low light environment can be challenging, for many reasons. Short exposure images tend to suffer from noise, while long exposure images are suffering from blurriness. Several techniques have been introduced to help reducing noise and blurriness, but none of them were perfect. Our team has decided to reproduce the “Learning to See in the Dark” paper from 2018, which was a big step towards improving low light photography. We chose to optimize the algorithm by using a smaller, but hopefully faster fully connected neural network, with U-net architecture, featuring batch-normalization, and in addition we created another convolutional neural network capable of predicting each images’ amplification ratio, based on the image itself as an input.

Kivonat

Fényképezés gyenge fényviszonyok között több okból kifolyólag is akadályokba ütközhet. Rövidebb záridő esetén a fényképünk nagy eséllyel zajos lesz, hosszabb záridő esetén pedig elmosódás lesz megfigyelhető. Az idők során több technikai újítás is érkezett, amik ezt a problémát voltak hivatottak orvosolni, de nem nagy sikerrel. A csapatunk egy 2018-ban kijött, “Learning to See in the Dark” című értekezés - ami megjelenésekor nagy lépés volt rossz fényviszonyok közötti fényképezés javítására - átalakítását tűzte ki célul. Tervünk, az előbb említett projekt algoritmusának optimalizálása egy kisebb, de remélhetőleg gyorsabb neurális hálózattal, ami a U-net architektúra alapjaira épül. Továbbá egy másik, konvolúciós neurális hálózatot is létrehoztunk, amely képes egy rossz fényviszonyok között készült képről - mint bemenetről - megállapítani annak amplifikációs értéket(amplification ratio).

Introduction

Imaging in low light circumstances is challenging. Using a higher ISO brightens the picture, however introduces a lot of noise, and using a wider aperture is often not possible since it is hardware constrained. The team in 2018 created a freely available dataset, consisting of more than 5000 image. Since these images are raw, the size of the compressed dataset is over 75 GB. There were two different datasets, Sony and Fuji, so we have chosen to go with the Sony dataset, which is more than 60 GB uncompressed. The images are not really interesting, they were taken in low light environments, with the exposure being between 1/10 and 1/30 of a second. For every low light image there is a given ground truth image(to be exact, for every 12 low light images there is a ground truth one) with an exposure between 10 and 30 sec. Meaning a 100 to 300 longer exposure(or 7 to 9 steps). The images are 4240×2832 pixels, which means passing through a whole image would be very memory intensive. Our team decided to use smaller patches and pass those through the network, leading to big improvements in time.

Related work

The idea of the project came from a paper published in 2018, called “Learning to See in the Dark”, done by four computer scientists. Their paper was a big step towards improving low light photography. After a few months of their release, Google came out with Night Mode, and several manufacturers followed. The Dataset they used - as previously mentioned - combined of a Sony and a Fuji dataset. Their network used two different fully convolutional network: the Sony dataset featuring CAN(mainly used for fast image processing), and the Fuji dataset featuring U-net. They used patches, sized 512x512, and performed random flipping for data enrichment. The paper leaves room for improvement. One of the problems of their model is that the amplification ratio must be chosen externally. The amplification ratio scales the brightness of the image to a desired value. So it would be useful to infer it from the input image.

Environment

We are using Jupyter Notebook, since a lot of visualizations are being shown through the coding phase. Due to hardware(mostly GPU and RAM) limitations, we set up a Google Cloud instance, featuring 1 * Nvidia Tesla K80 GPU, 500 GB SSD, n1-standard-2 (2 vCPUs, 7.5 GB memory), Ubuntu 18.04 LTS, and the zone being europe-west1-b. We installed CUDA, which is necessary for general computing on GPUs. Setting up the GPU environment did not go as planned, but after re-setup, GPU was available.

Network

Our project features two different neural networks for each task, both layers are convolutional:

First, the network for predicting the correct amplification ratio of a given raw input image. This network is a basic fully convolutional network featuring 2d convolutional, maxpooling, one flatten and finally a dense layer. The goal is to “shrink” the 3d image into one number corresponding to the required output value. The input image fed into this pipeline has to be 64x64x3 pixels. We use MSE(Mean Squared Error) to measure loss, since we would like to know the difference between the required output value and the prediction.

Second, the network trained on the raw images. This network is based on the U-net architecture. U-net is a convolutional network, consisting of a contracting path and an expansive path. The contracting path is based on 2d convolutional and maxpooling layers, meaning it works like an encoder, reducing the width, height of the input image. The expansive path is a sequence of up-convolutions and concatenation with high resolution features from the contracting path, which leads to a high-resolution output image. We tried to somehow improve the speed of this network by adding Batch Normalization, and Dropout. About the hyperparameters, for optimizer we went with Adam, with a learning rate of 0.001

Implementation

Collecting the data, preparation

As above mentioned, our dataset consists of images taken by the group that wrote the reference paper. We have approximately 2500 images - each image with its raw and ground truth image pair. The resolutions are 4240×2832 , thus the dataset is more than 60 GBs. After unzipping, the images are divided into three directories, train, valid and test, in these directories each line contains four texts, the path to the input image, the path to its ground truth image, the ISO score, and aperture. Then the raw images are loaded into the correct 2-dimensional arrays with their ground truth pair.

The two networks require different data preparation, described below.

The convolutional network predicting the correct amplification ratio uses $64 \times 64 \times 3$ images as input images. Since the raw images have way higher resolution, we resized the input raw images to fit the network. (For this we created a script that is capable of resizing the images to way smaller resolutions, using the cv2 library) The network uses amplification ratio values as desired outputs, so after normalization we store these in a set as well. We created six arrays, an input and an output pair for train, validation, and test, which later are fed into the cnn.

Our U-net based network

Training

Training on the smaller fully convolutional network was way faster than the other network. After the preparation of the data, we used the prepared training and validation input, output to measure the strength of the network. We used size 16 for batches, and the training lasted for 100 epochs. For callback functions, we used early-stopping monitoring the validation loss, to reduce the chances of overfitting. In addition, ReduceLROnPlateau was used, to help reducing learning rate, if validation loss does not improve over a given number of epochs (patience). At first, the network could not produce efficient outputs, loss and validation loss could not improve but decrease, even after 100 epochs, predictions were improper. To improve, we tried hyperparameter optimization.

Evaluation, Hyperparameter optimization

The network, trained on the $64 \times 64 \times 3$ images - as mentioned - could not improve, so optimizing the hyperparameters were necessary. We created a script for this process; using Hyperas, a simple wrapper around hyperopt for fast prototyping around Keras models. Following the Hyperas conventions, while building the model, every hyperparam had different options to choose from, layer feature sizes, the rate of dropout, type of optimizer, activation function, the value of learning rate, etc. At first, we tried training on 25 epochs for

one hyperparam set, and this method over 50 evaluations. The optimization lasted for approximately 30 minutes, so even higher number of evals is in an acceptable range. The best validation loss we achieved was 0.447, meaning with a small optimization, improvements were done.

Test

To start, we tested the smaller network's efficiency. The test dataset was prepared as the train and validation set, test input resized, output rescaled. The testing was done after optimizing the hyperparameters. The results were quite close to their ground truth pair, to be even more exact, we calculated their $(\text{abs}(\log_2(\text{prediction}) - \log_2(\text{truth})))$ value to get a better picture; here are some examples:

Predicted ratio	True ratio	Divergence(equation above)
262.92877	300	0.1902904777798966
270.19342	300	0.15096995728336893
245.1514	250	0.028255107865287066

We suppose the small difference between the ratios is due to the resizing of the images, since a lot of information got lost in that process.

Android App

To make the predictions more sightful, we came up with an Android Application, which is capable of - through user interaction - showing how the RAW images improve over the training. The main logic of the application is that it lets the user choose from different raw images at first, and after the evaluation it displays how the RAW image improved.

Future plans, Summary

...

One promising plan is to make our application capable of taking the actual raw images and then improving them, not just improving the previously downloaded raw images.

References

1. https://cchen156.github.io/paper/18CVPR_SID.pdf
2. <https://arxiv.org/pdf/1709.00643.pdf>

